

Analyze A/B Test Results

This project will assure you have mastered the subjects covered in the statistics lessons. The hope is to have this project be as comprehensive of these topics as possible. Good luck!

Table of Contents

- [Introduction](#)
- [Part I - Probability](#)
- [Part II - A/B Test](#)
- [Part III - Regression](#)

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question. The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the [RUBRIC](#).

Part I - Probability

To get started, let's import our libraries.

IMPORT AND VIEW DATA

```
[1]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

+ Code + Markdown

1. Now, read in the `ab_data.csv` data. Store it in `df`. Use your dataframe to answer the questions in Quiz 1 of the classroom.

a. Read in the dataset and take a look at the top few rows here:

```
[2]: df=pd.read_csv('../input/ab-test/ab_data.csv')
df.head()
```

Out[2]:

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.558738	control	old_page	0
1	804228	2017-01-12 08:01:45.158738	control	old_page	0
2	881590	2017-01-11 18:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	884975	2017-01-21 01:52:26.210827	control	old_page	1

CONTROL DATA – (MISSING ROWS, SHAPE ETC.)

b. Use the below cell to find the number of rows in the dataset.

```
[3]: df.shape[0]
```

```
Out[3]: 294478
```

c. The number of unique users in the dataset.

```
[1]: df['user_id'].nunique()
```

d. The proportion of users converted.

```
[4]: df['converted'].mean()
```

```
Out[4]: 0.11965919355685512
```

e. The number of times the `new_page` and `treatment` don't line up.

```
> df.groupby(['group', 'landing_page']).count()
#control group should have old page but there are 1928 people in control group who received new page
#treatment group should have new page but there are 1965 people in control group who received old page
```

```
Out[5]:
```

		user_id	timestamp	converted
group	landing_page			
control	new_page	1928	1928	1928
	old_page	145274	145274	145274
treatment	new_page	145311	145311	145311
	old_page	1965	1965	1965

```
[6]: df.query('group == "control" and landing_page == "new_page").nunique()+df.query('group == "treatment" and landing_page == "old_page").nunique()
#total 3893 people received incorrect page
```

```
Out[6]:
```

```
user_id      3893
timestamp    3893
group         2
landing_page  2
converted     4
dtype: int64
```

f. Do any of the rows have missing values?

```
[7]: df.isnull().sum()
#There is not missing value for columns
```

```
Out[7]:
```

```
user_id      0
timestamp    0
group         0
landing_page  0
converted     0
dtype: int64
```

CREATE DF2 WITH CORRECT VALUES

2. For the rows where treatment is not aligned with new_page or control is not aligned with old_page, we cannot be sure if this row truly received the new or old page. Use Quiz 2 in the classroom to provide how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in df2.

```
[8]: # I kept where group is treatment and page is new + group is control and page is old under df2 dataset
df2 = df[(df['group'] == 'treatment') == (df['landing_page'] == 'new_page')]
df2.head()
```

```
Out[8]:
```

	user_id	timestamp	group	landing_page	converted
0	881104	2017-01-21 22:11:48.558739	control	old_page	0
1	804228	2017-01-12 08:01:45.156739	control	old_page	0
2	881590	2017-01-11 18:55:08.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	884975	2017-01-21 01:52:28.210827	control	old_page	1

```
[9]: # Double Check all of the correct rows were removed - this should be 0
df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].shape[0]
```

```
Out[9]:
```

FIND AND DELETE DUPLICATED ROWS

3. Use df2 and the cells below to answer questions for Quiz3 in the classroom.

a. How many unique user_ids are in df2?

```
[10]: df2['user_id'].nunique()
```

```
Out[10]:
```

b. There is one user_id repeated in df2. What is it?

```
[11]: #there is 1 user_id repated in df2
df2['user_id'].duplicated().sum()
```

```
Out[11]:
```

```
[12]: print(df2[df2['user_id'].duplicated()])
#user_id 773192 is repeated at line 2893
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.598927	treatment	new_page	0

```
[13]: print(df2[df2['user_id'].duplicated()])
```

	user_id	timestamp	group	landing_page	converted
2893	773192	2017-01-14 02:55:59.598927	treatment	new_page	0

d. Remove one of the rows with a duplicate user_id, but keep your dataframe as df2.

```
[14]: #drop duplicates and save dataset to df2
df2=df2.drop_duplicates(subset=['user_id'])
```

```
[15]: #control if 2893.line is deleted
df2.query('user_id!="773192"')
```

```
Out[15]:
```

	user_id	timestamp	group	landing_page	converted
1899	773192	2017-01-09 05:37:58.781808	treatment	new_page	0

CALCULATE CONVERSION RATES OF OVERALL, CONTROL GROUP AND TREATMENT GROUP

4. Use df2 in the below cells to answer the quiz questions related to Quiz 4 in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

+ Code + Markdown

```
[16]: #converted percentage
df2['converted'].mean()
```

```
Out[16]: 0.11959708724499628
```

b. Given that an individual was in the control group, what is the probability they converted?

```
[17]: df2.query('group == "control"')['converted'].mean()
```

```
Out[17]: 0.1203863845004612
```

c. Given that an individual was in the treatment group, what is the probability they converted?

```
[18]: df2.query('group == "treatment"')['converted'].mean()
```

```
Out[18]: 0.11880806551510564
```

d. What is the probability that an individual received the new page?

+ Code + Markdown

```
▶ df2.groupby(['landing_page']).count()
#there are 145310 people who received new page
```

```
Out[19]:
```

	user_id	timestamp	group	converted
landing_page				
new_page	145310	145310	145310	145310
old_page	145274	145274	145274	145274

```
[20]: #percentage of people who received new page
#formula: new page / new page + old page
df2.query('landing_page=="new_page"').count()/len(df2)
```

```
Out[20]: user_id      0.500062
timestamp  0.500062
group      0.500062
landing_page 0.500062
converted  0.500062
dtype: float64
```

e. Consider your results from a. through d. above, and explain below whether you think there is sufficient evidence to say that the new treatment page leads to more conversions.

Your answer goes here.

My comment:

overall conversation: 0.1196 control group conversation: 0.1203 (%7 more than overall) treatment group conversation: 0.1189 (8% less than overall)

8% is small percentage. I can not say that new page causes more conversions. Old page seems better than new page at this line. We should do more tests to understand better if new page is really unsuccessful.

```
[21]: df2[df2['group']=='treatment'].timestamp.max(), df2[df2['group']=='treatment'].timestamp.min()
```

```
Out[21]: ('2017-01-24 13:41:44.097174', '2017-01-02 13:42:05.378582')
```

My comment:

Test is run for 22 days. This might be short time for users to understand if website is better or not.

DEFINE NULL AND ALTERNATIVE HYPOTHESES

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your *hypothesis in terms of words or in terms of p_{old} and p_{new}* , which are the converted rates for the old and new pages.

+ Code

+ Markdown

Put your answer here.

My comment:

My null hypotheses: H_0 = new page is unsuccessful or as successful as old page
My alternative hypotheses: H_1 = new page is more successful than old page

$H_0: p_{new} - p_{old} \leq 0$

$H_1: p_{new} - p_{old} > 0$

CALCULATE SUCCESS RATES

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the converted success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the converted rate in `ab_data.csv` regardless of the page.

Use a sample size for each page equal to the ones in `ab_data.csv`.

Perform the sampling distribution for the difference in converted between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use Quiz 5 in the classroom to make sure you are on the right track.

+ Code

+ Markdown

a. What is the convert rate for p_{new} under the null?

```
[22]: p_new = df2.converted.mean()
      p_new
```

Out[22] 0.11959708724499628

b. What is the convert rate for p_{old} under the null?

```
[23]: p_old = df2.converted.mean()
      p_old
```

Out[23] 0.11959708724499628

c. What is μ_{new} ?

c. What is n_{new} ?

```
[24]: n_new = df2.query('landing_page == "new_page"')['converted'].count()
      n_new
```

Out[24]: 145310

d. What is n_{old} ?

```
[25]: n_old = df2.query('landing_page == "old_page"')['converted'].count()
      n_old
```

Out[25]: 145274

SIMULATE SUCCESS RATES WITH BINOMIAL FUNCTION

e. Simulate n_{new} transactions with a convert rate of p_{new} under the null. Store these n_{new} 1's and 0's in new_page_converted.

```
[26]: #binomial function draws samples of normal distribution
      # 1 trial, probability is p_new(0.1195) and we will do this n_new (145310) times
      new_page_converted = np.random.binomial(1, p_new, n_new)
```

new_page_converted

Out[27]: array([0, 0, 0, ..., 0, 0, 0])

```
[28]: new_page_converted.mean()
```

Out[28]: 0.11908333906819983

f. Simulate n_{old} transactions with a convert rate of p_{old} under the null. Store these n_{old} 1's and 0's in old_page_converted.

```
[29]: old_page_converted = np.random.binomial(1, p_old, n_old)
```

```
[30]: old_page_converted.mean()
```

Out[30]: 0.11889945895342593

g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).

```
[54]: obs_simulation_diff = new_page_converted.mean() - old_page_converted.mean()
      obs_simulation_diff.mean()
```

Out[54]: 0.00018388011477309119

+ Code

+ Markdown

My comment:

observed differences convert rate of pages is not significant. But we should evaluate it 10000 times with bootstrapping and observe difference again.

CALCULATE SUCCESS RATES OF TREATMENT AND CONTROL GROUP

h. Simulate 10,000 $p_{new} - p_{old}$ values using this same process similarly to the one you calculated in parts a. through g. above. Store all 10,000 values in a numpy array called `p_diffs`.

+ Code

+ Markdown

```
[32]: df2.head()
```

```
Out[32]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.550739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	081590	2017-01-11 18:55:08.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143785	treatment	new_page	0
4	884975	2017-01-21 01:52:28.210827	control	old_page	1

```
[88]: #diffs of control team converted rate and treatment team converted rate
obs_diff = df2.query('group == "treatment"')['converted'].mean() - df2.query('group == "control"')['converted'].mean()
obs_diff
```

```
Out[88]: -0.001578238985355567
```

My comment:

Observed differences is -0.00157. Really small. Now I will evaluate differences with bootstrapping method

SIMULATE SUCCESS RATES OF GROUPS WITH BOOTSTRAPING

```
[48]: # create sampling distribution of difference in average converted rate
# with bootstrapping
diffs = []
size = df2.shape[0]

for _ in range(1000):
    b_samp = df2.sample(size, replace=True)
    control_mean = b_samp.query('group == "control"').converted.mean()
    treatment_mean = b_samp.query('group == "treatment"').converted.mean()
    diffs.append(treatment_mean - control_mean)
```

+ Code

+ Markdown

My note:


I could not arrange my range to 10000. My cpu can not compute it.

My comment:

I appended every differences to diffs array for each samples.

DRAWING HISTOGRAM

```
diffs = np.array(diffs)
# 99% confidence interval
low = np.percentile(diffs, .5)
upper = np.percentile(diffs, 99.5)
plt.hist(diffs);
plt.axvline(x=low, color='red', linewidth=2);
plt.axvline(x=upper, color='red', linewidth=2);
plt.axvline(obs_diff, color='yellow', linewidth=2);
```



+ Code

+ Markdown


My comment:

I can see bell shaped graphic, so there is normal distribution. confidence interval is %99 and almost most of samples' mean are between lower and upper line

```
null_vals = np.random.normal(0, diffs.std(), diffs.size)
plt.hist(null_vals)

# plot line for observed statistic
plt.axvline(obs_diff, c='red')
```

Out[69]: <matplotlib.lines.Line2D at 0x7f93b8f36a50>



j. What proportion of the p_diffs are greater than the actual difference observed in ab_data.csv?

```
[79]: #p value calculation
      (null_vals > obs_diff).mean()
```

Out[79]: 0.919

COMMENT

My comment:

Type I error is 0.5

P value is greater than 0.5.

**** We can not reject null hypothesis ****